



The vNext of Intelligent Observability

By José Eserich

Enterprise Architect | Modernization Strategist | vNext Framework Author

2025

From Chaos to Clarity. From Telemetry to Understanding. From Monitoring to Autonomous Improvement.

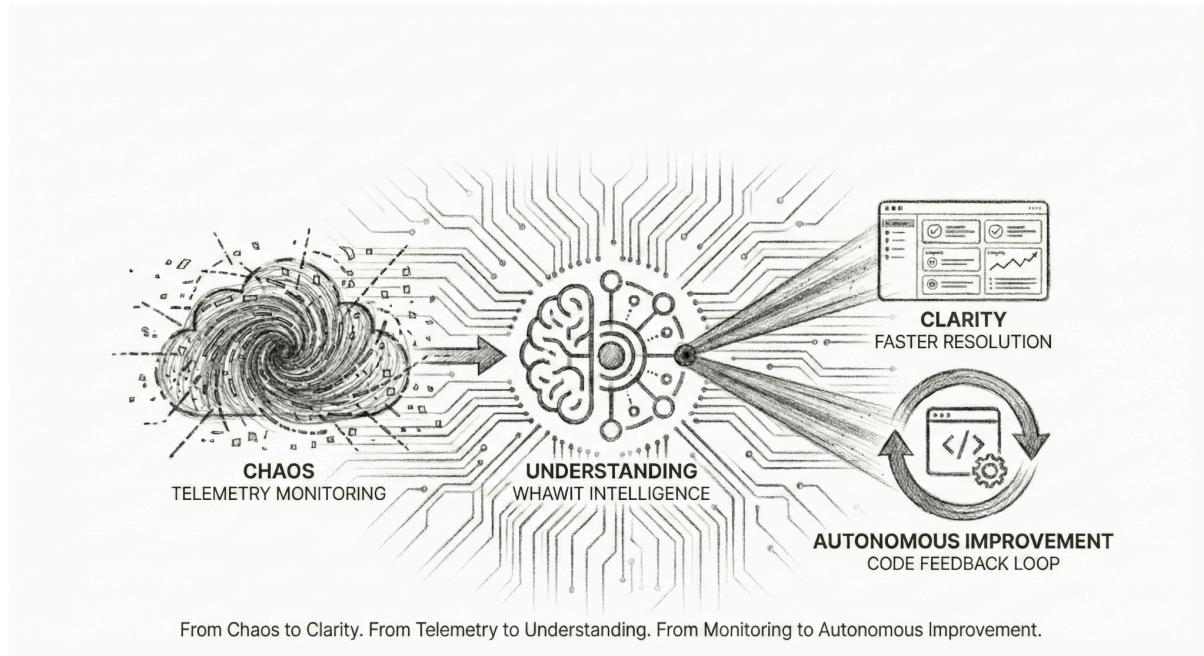
1. Summary

Modern software systems generate enormous volumes of telemetry—logs, metrics, events, traces—yet teams still depend on human interpretation to understand what any of it means. Tools have evolved to collect more data, display more dashboards, and trigger more alerts, but the core problem remains: **observability, as practiced today, rarely explains anything by itself.**

This has a measurable economic cost. Industry studies estimate that IT downtime can cost from **hundreds of thousands up to millions of dollars per hour**, depending on company size and sector. Mean Time To Recovery (MTTR) becomes not just a technical metric, but a direct financial KPI. At the same time, organizations are spending aggressively on observability platforms and log storage, often retaining terabytes of data that are rarely accessed unless something goes wrong.

WHAWIT introduces a new category: **Intelligent Observability**. It transforms raw telemetry into structured understanding by interpreting logs, correlating events, identifying root causes, and accelerating incident response. It is designed to reduce investigation time dramatically, thereby lowering MTTR, saving engineering hours, and reducing the real cost of downtime.

Beyond that, WHAWIT closes the loop. By analyzing code repositories and understanding how systems evolve over time, it can map production behavior back to the codebase, propose improvements, and even generate candidate fixes. Observability turns from a passive reporting layer into an active, autonomous reliability partner.



2. Context and Problem Definition

The move toward distributed, cloud-native architectures—microservices, containers, serverless, event-driven systems—has multiplied both the **volume** of telemetry and the **complexity** of incidents. A single customer-facing failure can involve dozens of microservices, multiple data stores, and several external providers.

When a major incident occurs, every minute matters. For many organizations, a single hour of downtime can exceed **\$300,000**, and for large enterprises it can reach or surpass **seven figures per hour**. In that window, engineering teams are typically:

- Manually correlating logs across services.
- Reconstructing the timeline of events by hand.
- Guessing which recent code changes might be responsible.

Meanwhile, the cost continues to accumulate.

Paradoxically, these same organizations often maintain huge log archives that are rarely used. Terabytes of logs are collected “just in case,” but **only a small fraction is ever queried** outside of incident windows. Log volumes grow year over year, storage and licensing costs escalate, and many teams react by reducing retention or sampling aggressively, increasing the risk of flying blind when a real incident occurs.

In this environment, the challenge is no longer data collection. The challenge is **turning that data into fast, reliable understanding without exploding cost and complexity**.

3. Why Traditional Observability Falls Short

Traditional observability stacks are excellent at collecting and visualizing data, but they stop short of genuine understanding.

Dashboards and metrics can show that an error rate spiked at 10:03, or that latency jumped on a specific endpoint, but they do not explain why. Powerful log search tools allow engineers to query millions of lines of logs, but finding the right subset—and understanding the narrative behind it—remains a manual, time-consuming task.

This model has three structural problems:

1. **Time to Insight:** When incidents happen, engineers lose tens of minutes—or hours—sifting through logs. This directly inflates MTTR and the cost of downtime.
2. **Engineering Time Waste:** Studies consistently show that a significant portion of engineering effort (sometimes 20–30% of working time) is consumed by unplanned work: incident response, firefighting, and manual diagnostics.

This paper was **written entirely by the author without the assistance of generative AI**.

Some illustrative images associated with the publication may have been **generated using AI-based tools** for conceptual visualization purposes only.
© 2025 — All rights reserved. No part of this document may be reproduced, distributed, or transmitted without the explicit permission of the author.

3. **Unmanaged Telemetry Economics:** Log data grows exponentially, while budgets do not. Many organizations see observability and log storage costs increasing year over year, often with surprise overage bills. To cope, they limit what they collect or how long they retain it, introducing blind spots.

The result is an unsatisfying trade-off: either accept ever-growing costs, or accept degraded visibility when it matters most.

4. Vision: Observability as Understanding

WHAWIT is built on a different premise: **observability should culminate in understanding and improvement, not just data exposure.**

Instead of focusing on more dashboards or more queries, WHAWIT focuses on **semantic interpretation**:

- Understanding what the logs are *saying*, not just storing them.
- Identifying which patterns actually matter in the context of an incident.
- Explaining incidents in clear language, with evidence.

The vision is to give teams immediate clarity when something breaks, cutting MTTR by a large factor and shrinking both downtime cost and engineering toil. At the same time, WHAWIT is designed to make log storage more economically rational by emphasizing **signal over raw data volume**, without sacrificing incident readiness.

5. Architecture Overview

WHAWIT's architecture is organized into several coordinated layers:

- **Ingestion Layer** – Connects to existing observability ecosystems (Datadog, New Relic, CloudWatch, Elastic, Kubernetes, etc.) and ingests logs regardless of their format or origin. WHAWIT does not attempt to replace these platforms; it **sits on top of them and amplifies their value**.
- **Semantic Processing Layer** – Normalizes and structures logs and events so they can be interpreted in a consistent way, even when they come from heterogeneous systems.
- **Intelligence Layer** – Interprets logs using context-aware analysis to identify anomalies, failure signatures, correlations across microservices, and probable root causes.

This paper was **written entirely by the author without the assistance of generative AI**.

Some illustrative images associated with the publication may have been **generated using AI-based tools** for conceptual visualization purposes only.
© 2025 — All rights reserved. No part of this document may be reproduced, distributed, or transmitted without the explicit permission of the author.

- **On-Call Hub** – Presents a real-time narrative of the incident: what is happening, why it is happening, and where to look first. It supports collaboration, escalation, and handovers.
- **Repository & CI Integration** – Connects with Git repositories and CI/CD pipelines to understand how the codebase evolves, which areas are fragile, and how incidents map back into the code.

Together, these layers transform the telemetry firehose into a coherent, actionable incident story—with a direct line into the code that produced it.

6. The Intelligence Layer

At the core of WHAWIT is an intelligence layer that **reads logs more like a senior engineer than a database**.

Rather than matching fixed patterns or relying solely on thresholds, it builds a contextual understanding of what is normal for a system and what is not. It can:

- Distinguish between routine noise and meaningful anomalies.
- Correlate events across services and infrastructure layers.
- Extract the small fraction of log lines that explain a failure.
- Summarize an incident in a concise, readable way.

The outcome is a significant reduction in the time it takes to move from “something is wrong” to “we know roughly what is happening and where to look.” That reduction translates directly into lower MTTR and fewer engineering hours lost to manual investigation.

7. Incident Acceleration and MTTR Reduction

Every minute shaved off MTTR reduces the financial and operational impact of an outage. For many organizations, even a 20–30% reduction in MTTR can save **hundreds of thousands of dollars per year** in avoided downtime; for large enterprises, the potential savings are much higher.

WHAWIT addresses MTTR in three stages:

1. **Faster Understanding:** Instead of opening a blank search bar, the on-call engineer opens WHAWIT and sees an interpreted narrative of the incident, with the most relevant evidence surfaced automatically.
2. **Reduced Engineering Toil:** Because the platform performs much of the initial triage and correlation, engineers spend less time hunting and more time fixing. Over a year, this can reclaim thousands of engineering hours across a team.

This paper was **written entirely by the author without the assistance of generative AI**.

Some illustrative images associated with the publication may have been **generated using AI-based tools** for conceptual visualization purposes only.
© 2025 — All rights reserved. No part of this document may be reproduced, distributed, or transmitted without the explicit permission of the author.

3. **Better Post-Incident Learning:** WHAWIT captures incident narratives and connects them to code changes, creating a richer knowledge base that makes future incidents easier to resolve.

The business result is a combination of **lower direct downtime cost** and **higher effective engineering capacity**.

8. The On-Call Hub

The On-Call Hub is the operational face of WHAWIT. It is designed around the reality that **most observability data is used under stress**—in the middle of an incident.

Instead of a collection of disconnected dashboards, the Hub provides:

- A timeline of key events.
- A synthesized explanation of what is happening.
- Links into relevant logs, metrics, and code.
- Context for handovers and escalations.

This reduces cognitive load for responders, supports better coordination, and shortens the “chaotic phase” of incident response.

9. Security and Compliance

WHAWIT is built with enterprise security and compliance requirements in mind:

- Encryption in transit and at rest.
- Strong tenant isolation.
- Support for zero-retention or short-retention modes where regulations demand it.
- Audit trails for incident access and actions.

The platform can operate in environments where telemetry data is sensitive, while still providing the intelligence layer that teams require.

10. Enterprise Ecosystem and Integrations

WHAWIT does not attempt to replace existing observability platforms. Instead, it **fuses their data into a higher-level understanding**.

It integrates with:

- Monitoring and APM tools such as Datadog, New Relic, and Elastic.

This paper was **written entirely by the author without the assistance of generative AI**.

Some illustrative images associated with the publication may have been **generated using AI-based tools** for conceptual visualization purposes only.
© 2025 — All rights reserved. No part of this document may be reproduced, distributed, or transmitted without the explicit permission of the author.

- Cloud-native logs from platforms such as AWS CloudWatch and Kubernetes.
- Collaboration tools like Slack, Teams, and PagerDuty.
- Git hosting platforms and CI/CD systems.

For both startups and enterprises, this means WHAWIT can be adopted incrementally, riding on top of existing investments and amplifying their ROI rather than displacing them.

11. The Autonomous Code Feedback Loop

A defining capability of WHAWIT is its ability to **close the loop** between what happens in production and how the codebase evolves.

The platform continuously analyzes repositories, commit histories, and architectural patterns. When incidents occur, WHAWIT relates them to specific components or recent changes in the code. It understands, for example, which services are historically fragile under load, which modules have a high concentration of failures, and which commits preceded a spike in incidents.

From that understanding, WHAWIT can propose concrete improvements: better error handling paths, resilience patterns, adjustments to configuration, or enhancements to test coverage. In many cases, it can generate candidate patches or refactorings and submit them as pull requests for human review.

Over time, as fixes are merged and the system stabilizes, WHAWIT learns which interventions are effective. The result is a feedback loop where incidents do not just get resolved; they **teach the system how to improve the code that caused them**.

This turns observability into autonomous reliability engineering, where the same platform that explains failures also helps prevent their recurrence.

12. Business Value and ROI

From a business perspective, WHAWIT addresses four intertwined cost centers:

1. **Downtime Cost:** By reducing MTTR, WHAWIT directly cuts the financial impact of outages, which can reach hundreds of thousands or millions per hour for serious incidents.
2. **Engineering Time:** By automating triage and correlation, WHAWIT can give back a significant portion of the 20–30% of engineering time often lost to reactive incident work.
3. **Log Storage Economics:** By focusing on high-value signals and intelligent interpretation, WHAWIT helps organizations make better use of existing observability investments, rather than simply ingesting and storing more data indefinitely.

This paper was **written entirely by the author without the assistance of generative AI**.

Some illustrative images associated with the publication may have been **generated using AI-based tools** for conceptual visualization purposes only.
© 2025 — All rights reserved. No part of this document may be reproduced, distributed, or transmitted without the explicit permission of the author.

4. **Code Quality and Resilience:** Through the autonomous feedback loop, incident learnings are translated into real code improvements, compounding the reliability gains over time.

The combination is a strong ROI story: **less downtime, less waste, and more value from data and people you already pay for.**

13. Roadmap

WHAWIT's roadmap deepens the same principles that define the platform today:

- Richer predictive capabilities to anticipate incidents before they fully manifest.
- More advanced, context-aware remediation suggestions.
- Stronger models for learning from incident history and code evolution.
- Tighter integration with CI/CD to enable safer, faster iteration.

The long-term direction is clear: systems that not only understand their own failures, but increasingly **anticipate and correct them proactively.**

WHAWIT represents a shift from traditional monitoring—collecting and presenting telemetry—to **intelligent, autonomous understanding and improvement.**

It accepts the realities of modern software: high complexity, high data volume, high cost of downtime, and finite engineering capacity. Then it redefines observability around what truly matters: faster understanding, better decisions, and a closed loop between production behavior and the code that powers it.

In doing so, WHAWIT turns observability from a cost center into a strategic asset.